

Preliminary Evaluation of MeshEcho Confidence-Aware Routing for Airtime-Constrained LoRa Mesh IoT Networks

Anonymous Authors

Paper submitted to ICCT 2026 Track 8: Internet-of-Things & Sensor Networks

Abstract—LoRa mesh networking is attractive for low-power Internet-of-Things (IoT) deployments that need wider coverage than a single hop can provide. However, multi-hop LoRa routing has an inherent reliability-airtime tension: redundant flooding improves reachability but increases channel occupancy, while cached source routes reduce transmissions but can fail when a selected path contains a weak or stale hop. This paper presents a preliminary packet-level evaluation of that tradeoff for small LoRa mesh IoT networks. We build a reproducible simulator with a LoRa time-on-air model, log-distance path loss, log-normal shadowing, half-duplex radios, collision/capture behavior, and probabilistic reception from SNR margin. Using this harness, we compare managed flooding, a source-route cache baseline, and a lightweight confidence-aware routing prototype, MeshEcho, that ranks discovered paths and uses bounded recovery forwarding for route-recovery cases. In the frozen 50-node mixed-traffic matrix, MeshEcho reaches 0.735 ACK-confirmed unicast PDR versus 0.389 for the source-route baseline, with 839.6 s versus 971.5 s of aggregate airtime. A separate controlled route-conflict experiment shows that the route-selection gain appears only when the shorter candidate contains a materially weaker hop: with 7.5 dB extra loss on that hop, MeshEcho reaches 0.999 ACK-confirmed PDR versus 0.784 for a shortest-path ablation, whereas null, moderate, and reverse controls produce no gain. The evaluation also exposes important fairness boundaries: fixed repeated pairs favor route caching, the default link setting is link-quality saturated and produces very few multi-hop routes, and adaptive recovery is not budget-matched across all protocol variants. These results support a bounded preliminary claim rather than universal dominance.

Index Terms—LoRa mesh, Internet of Things, sensor networks, adaptive routing, airtime efficiency, packet delivery ratio

I. Introduction

LoRa radios are widely used in low-power IoT systems because they provide long communication range at modest device cost and energy consumption [1]. Recent surveys and overview articles show that multihop LoRa/LoRaWAN networking can relax star-topology constraints, but it must manage airtime, link variability, and control overhead carefully [2]–[5]. In many deployments, terrain, indoor attenuation, transmit-power limits, and gateway placement constraints can leave coverage gaps. LoRa mesh networking addresses this by allowing intermediate devices to relay packets across multiple hops for environmental sensing, temporary field monitoring, emergency communication, and community-scale telemetry.

The challenge is that LoRa’s low data rate and long packet time on air make redundant forwarding expensive. A flooding-oriented design can reach many nodes, but every rebroadcast consumes scarce airtime and raises collision pressure. A source-route cache can reduce repeated transmissions, but a cached route is only useful while its weakest hop remains good. This creates a practical design question for LoRa mesh IoT networks:

How much forwarding redundancy should a node spend when path quality is uncertain and airtime is scarce?

Recent protocol studies explore several points in this design space, including distance-vector LoRa mesh routing, reusable mesh libraries, spreading-factor clustering, low-latency multihop forwarding, and real-time industrial LoRa communication [6]–[10]. This paper studies the question as a preliminary routing-system evaluation. We compare two fixed-mode baselines with a lightweight confidence-aware protocol named MeshEcho. MeshEcho is the name of our protocol and is not a fork of Meshtastic or MeshCore. The prototype is intentionally conservative: it does not claim to solve all adaptive LoRa mesh routing problems. Instead, it evaluates whether modest path confidence information can expose a better operating point between managed flooding and pure source-route caching.

The contributions of this paper are:

- We present a reproducible packet-level LoRa mesh simulation harness that evaluates routing policies under common topology, traffic, PHY, propagation, collision, and airtime assumptions.
- We characterize the reliability-airtime tradeoff between managed flooding and cached source routing in repeated-unicast and mixed IoT traffic.
- We evaluate MeshEcho, a preliminary confidence-aware routing protocol that ranks candidate routes using local link-quality evidence and limits fallback flooding to bounded recovery cases.
- We separate general random-topology results from a controlled route-conflict experiment with null and reverse controls, making the mechanism claim and its fairness boundary explicit.

The scope is deliberately preliminary. We focus on simulation-based evidence and small-scale protocol charac-

terization for ICCT Track 8. The implementation reports the reproducible calm-mesh line as MeshEcho; the later online Smart-CALM controller is treated as an exploratory extension rather than as the main claim of this manuscript. A more complete firmware-oriented controller and broader system validation are left for future work.

II. Background and Motivation

A. LoRa Mesh Routing Tradeoffs

LoRa modulation trades data rate for link budget. For IoT sensing, this is often acceptable because payloads are small and traffic is intermittent. In a mesh, however, each relay repeats the same low-rate packet. A route with several hops therefore consumes multiple packet airtimes, and a flood can multiply this cost across many relays.

Current LoRa mesh systems often operate near one of two routing extremes. A managed-flooding approach rebroadcasts packets with randomized delays and duplicate suppression. This is robust because delivery does not depend on one chosen path, but it can generate many transmissions. A source-route cache first discovers a path and then reuses it for later unicast packets. This reduces redundant forwarding when traffic pairs repeat, but can be brittle if the first discovered path includes a weak hop.

The same tension appears in low-power and lossy network routing more broadly. RPL uses objective functions and routing metrics for constrained networks [11], while ETX-style metrics showed that path selection should consider delivery probability rather than hop count alone [12]. LoRa mesh differs in that airtime is particularly expensive and packets are long enough for collision pressure to dominate at modest traffic levels. This motivates a routing policy that treats redundancy as a limited resource rather than as a fixed protocol behavior.

B. Design Goal

Our design goal is not to maximize reliability at all costs. Instead, we seek a policy that can:

- 1) keep the efficiency benefit of source routing when a path appears sufficiently reliable,
- 2) avoid accepting obviously weak discovered paths when alternatives are available, and
- 3) use bounded fallback forwarding only when route discovery or route confidence indicates uncertainty.

This framing is suitable for IoT and sensor-network settings where traffic is light enough for multi-hop LoRa to be useful, but dense enough that unnecessary rebroadcasts still reduce network capacity.

III. System Model

A. Network and Radio Model

We simulate N LoRa nodes uniformly placed in a square area. Nodes share one half-duplex LoRa channel and may relay packets unless explicitly configured otherwise. The default radio profile uses spreading factor SF9, bandwidth 125 kHz, coding-rate index 1 (4/5), 32-byte payloads,

17 dBm transmit power, and a capture threshold of 6 dB, matching a representative SX126x-class LoRa radio configuration [13].

Packet time on air follows the standard LoRa symbol-time calculation [14]. For spreading factor SF , bandwidth BW , payload length PL , coding-rate index CR , explicit header flag IH , CRC flag CRC , and low-data-rate optimization flag DE , the symbol time is

$$T_{sym} = \frac{2^{SF}}{BW}. \quad (1)$$

The packet time on air is the preamble time plus the payload symbol time:

$$T_{oa} = (N_{pre} + 4.25)T_{sym} + N_{payload}T_{sym}, \quad (2)$$

where $N_{payload}$ is computed using the standard LoRa payload-symbol expression. The simulator accumulates total airtime by adding T_{oa} for each transmitted packet.

Propagation uses log-distance path loss with log-normal shadowing:

$$PL(d) = PL_0 + 10n \log_{10}(d/d_0) + X_\sigma, \quad (3)$$

where n is the path-loss exponent and X_σ is a zero-mean Gaussian shadowing term. The received power is $P_r = P_t - PL(d)$. The receiver noise floor uses thermal noise plus a fixed noise figure. The simulator maps SNR margin to packet reception probability through a logistic function:

$$P_{rx} = \frac{1}{1 + \exp(-k\Delta_{SNR})}, \quad (4)$$

where Δ_{SNR} is the difference between measured SNR and the required LoRa SNR threshold for the configured spreading factor. When transmissions overlap, a receiver applies a capture rule: the desired packet can survive only if it is at least the capture threshold stronger than the strongest interferer. For tractability, all packet kinds use the same configured payload length in the current simulator. Source-route header growth and protocol-specific control-frame layouts are therefore not separately charged in the airtime model.

B. Traffic Model and Metrics

We evaluate repeated unicast traffic and mixed traffic. Repeated unicast uses a fixed set of source-destination pairs so that route caching has a meaningful opportunity to help. Mixed traffic contains both unicast and broadcast flows, which is closer to many IoT mesh deployments that combine telemetry, control, and local announcements.

The main metrics are unicast packet delivery ratio (PDR), broadcast coverage, total transmissions, total airtime, airtime per delivery, and collision failures. We report means over 20 random seeds. Each seed changes node placement, traffic timing, selected pairs, the static per-link shadowing realization, and stochastic reception outcomes. The frozen main matrix also uses the historical shared random stream in which forwarding jitter and reception

draws can consume the same generator. The released simulator provides an opt-in independent channel-randomness mode for new runs; this separates reception draws from forwarding jitter and learning exploration, but it still does not create event-by-event common random numbers when protocols generate different numbers of transmissions. In the main matrix, `unicast_pdr` is the source-side ACK-confirmed delivery ratio; destination DATA arrival is reported separately because destination reception does not guarantee a successful ACK return. The main tables use the frozen three-protocol matrix (managed flooding, source-route cache, and calm-mesh); the later Smart-CALM ablation files are retained as historical artifacts and are not used for the quantitative claims in this paper. The comparison is a common-harness comparison, not a claim that every protocol has identical internal recovery semantics: MeshEcho has bounded fallback behavior, whereas the two fixed baselines do not have an identical recovery controller.

IV. Routing Policies

A. Managed Flooding Baseline

The managed-flooding baseline models a Meshtastic-like forwarding behavior at a protocol level [15]. A source transmits a data packet. Relay-capable nodes that hear a new packet schedule a delayed rebroadcast. If a relay hears another copy of the same packet before its scheduled transmission, it suppresses its own forward. Packets stop when the destination receives them or when the hop limit expires. This baseline represents a high-redundancy operating point.

B. Source-Route Cache Baseline

The source-route baseline models a MeshCore-like path-cache behavior [16]. If a source has no cached route to a destination, it floods a route request (RREQ). The destination returns a route reply (RREP) along the reverse discovered path. The source then caches the path and later unicast packets carry the complete source route. Only the next hop on the source route forwards the packet. This baseline represents a low-redundancy operating point.

C. MeshEcho Confidence-Aware Prototype

MeshEcho keeps the same basic route-discovery structure but changes how route candidates are admitted. During a short discovery window, the destination may observe multiple RREQ paths. Instead of immediately accepting the first path, the prototype scores candidate paths using simple local evidence from received packets. At a high level, the score increases with stronger SNR margin and decreases with collision evidence and additional hops. The selected route is cached with a finite lifetime.

The prototype also includes bounded fallback forwarding. If no route is found, the source sends a fallback flood to recover reachability. The implementation also supports

a confidence-triggered delayed fallback behind a source-routed packet, which is canceled if the original delivery succeeds. In the frozen calm-mesh matrix, the confidence threshold is set to zero, so the reported CALM fallback counts are driven by route-recovery cases rather than by an independently tuned confidence-triggered flood. Fallback is therefore a small-radius recovery tool, not a permanent flooding mode.

This section intentionally describes only the preliminary MeshEcho mechanism evaluated in this paper. More aggressive adaptive recovery and implementation-specific control policies are outside the scope of this submission.

V. Experimental Evaluation

A. Setup

Unless stated otherwise, simulations use 50 nodes in a 3000 m $\times$ 3000 m square area, 1200 s duration, eight repeated unicast pairs, maximum hop count 7, SF9, 125 kHz bandwidth, 32-byte packets, 17 dBm transmit power, path-loss exponent 2.7, shadowing standard deviation 4 dB, and 20 random seeds. Mixed traffic uses a mean offered load of six application packets per minute. The high-shadowing stress case increases the shadowing standard deviation to 6 dB. These settings define a reproducible common-harness comparison, but they are not intended to represent every deployment density or radio range. The repeated-pair workload is deliberately retained as a route-cache efficiency case; it is not a neutral pair-churn workload. The reported tables retain this frozen legacy matrix for traceability rather than silently replacing it with a harder post-hoc parameter choice.

We performed a post-hoc calibration audit without replacing the frozen results. Across 24,500 unordered node pairs, the minimum calculated direct-link reception probability was 0.999909, with no pair below 0.99. The fixed-pair minimum was 0.9999996 and remained 0.9999866 after an additional 8 dB penalty. In the replayed main runs, 128 of 129 MeshCore cached routes and all 142 CALM cached routes were one hop. The matrix is therefore dominated by contention and control traffic rather than multi-hop link-quality conflicts. The static per-link shadowing realization is shared across protocols for each seed, but temporal fading is not modeled. New fairness-oriented reruns should enable the simulator's `-independent-rng-streams` option, while the frozen files remain in legacy mode for exact reproduction of the published numbers.

B. Repeated Unicast

Table I and Figs. 2-3 show the repeated-unicast result. Managed flooding reaches the highest unicast PDR in this scenario at 0.838, but spends 1152 s of aggregate airtime and produces 120,276 collision failures. Source-route caching reduces airtime to 358 s, but its PDR drops to 0.605. MeshEcho reaches 0.831 PDR with 256 s of airtime and 17,619 collision failures. Thus, MeshEcho is close to flooding reliability in this particular repeated-pair




Fig. 1. Implemented MeshEcho routing architecture. Route discovery collects candidate paths, receive evidence is converted to path confidence, and cached source routes are supplemented only by bounded fallback forwarding when path confidence or route availability is insufficient.

TABLE I
Repeated-unicast results over 20 seeds. MeshEcho denotes the confidence-aware prototype.

Protocol	Unicast PDR	Total airtime (s)	Collision failures	Airtime reduction vs. flooding
Managed flooding	0.838	1152	120276	—
Source-route cache	0.605	358	27488	68.9%
MeshEcho	0.831	256	17619	77.7%

TABLE II
Mixed-traffic and high-shadowing results over 20 seeds. MeshEcho denotes the confidence-aware prototype.

Scenario	Protocol	Unicast PDR	Broadcast coverage	Total airtime (s)	Collision failures
Mixed traffic	Managed flooding	0.852	0.947	1183	123501
Mixed traffic	Source-route cache	0.389	0.960	971	82454
Mixed traffic	MeshEcho	0.735	0.968	840	74103
High shadowing	Managed flooding	0.842	0.952	1175	123322
High shadowing	Source-route cache	0.387	0.962	970	82324
High shadowing	MeshEcho	0.731	0.965	834	74129

setting while retaining most of the source-route efficiency. Because nearly all cached paths in this matrix are one hop, this result should not be interpreted as a general multi-hop route-selection result.

This scenario favors route caching because the same source-destination pairs repeat. It is therefore a useful cache-efficiency case, but not a neutral workload for protocols that do not rely on route reuse. The result should be read as a tradeoff point rather than as evidence that MeshEcho dominates flooding or generalizes to arbitrary pair churn. A fair follow-up should report random-pair traffic separately and should pre-register a non-degenerate direct-link PRR distribution and multi-hop route fraction.

C. Mixed Traffic

Mixed traffic is more challenging because broadcast and unicast flows share the same channel. Table II shows that the source-route baseline maintains high broadcast coverage but has a lower ACK-confirmed unicast PDR of 0.389, compared with 0.852 for managed flooding. MeshEcho recovers part of this loss, reaching 0.735 while

reducing total airtime from 971 s to 840 s relative to the source-route baseline. Collision failures decrease from 82,454 to 74,103.

Compared with managed flooding, MeshEcho reduces airtime by 29.0% and collision failures by 40.0%, while keeping broadcast coverage within the same range. This end-to-end gain combines route admission and bounded recovery; it should not be attributed to confidence ranking alone. The remaining PDR gap to managed flooding indicates that the preliminary policy spends less redundancy but does not fully replace flood diversity.

D. High-Shadowing Stress Case

The high-shadowing case increases log-normal shadowing from 4 dB to 6 dB. The same qualitative pattern remains, but the absolute values are not identical to the mixed case. Source-route caching reaches 0.387 unicast PDR, while MeshEcho reaches 0.731. Total airtime decreases from 970 s to 834 s and collision failures decrease from 82,324 to 74,129. This shows robustness of the aggregate configuration under stronger shadowing, but



Fig. 2. Unicast delivery ratio across the three evaluation scenarios.



Fig. 3. Aggregate LoRa airtime across the three evaluation scenarios.

does not isolate confidence-aware admission because the random matrix still rarely creates a meaningful length-versus-quality conflict.

E. Controlled Route-Conflict Ablation

The random-topology matrix does not reliably create a direct conflict between shortest-path and confidence ranking, so we use a separate controlled experiment to test that mechanism. The topology has two candidate routes from node 0 to node 5: a short path 0–1–2–5 and a longer path 0–3–4–6–5. Non-edges are blocked explicitly. Each run uses SF9, zero stochastic shadowing, a 12 s discovery window, a 1200 s route-cache TTL, 24 unicast flows spaced by 30 s, and one route-discovery attempt. Online learning, timeout retry, and route-miss fallback are disabled. The confidence variant selects the highest-confidence candidate; the no-confidence variant selects the shortest candidate.




Fig. 4. Mixed-traffic reliability-airtime tradeoff under the evaluated protocol budgets. MeshEcho moves the source-route operating point toward higher delivery with lower airtime, but the end-to-end difference includes bounded recovery not available to the fixed baselines.

The main mechanism condition adds 7.5 dB path loss to link 1–2, while all other permitted links receive no additional loss. The fairness audit also contains a null condition, a moderate 5 dB short-branch condition, and a reverse condition that weakens link 4–6 on the long branch. Results are paired by seed and use ACK-confirmed PDR.

Table III shows that MeshEcho does not create a gain when the two branches are effectively equivalent, and it does not blindly prefer the longer route when the long branch is weakened. The 7.5 dB condition is therefore evidence that the candidate-ranking mechanism can detect a short-but-weaker path under this model. It is not evidence of topology-independent superiority. The higher airtime in that condition also shows the cost of selecting the reliable four-hop route instead of the three-hop route. The moderate +5 dB condition is an important negative control: the observed mean delta is zero, so the mechanism does not produce a gain for every small perturbation.

VI. Discussion

The preliminary results support four observations. First, fixed flooding and fixed source-route caching expose different sides of the LoRa mesh design space. Flooding protects delivery but consumes substantial airtime. Cached source routing saves airtime, but it can underperform when the selected route is fragile. Because the default direct links are almost all high-PRR, the random matrix should be interpreted primarily as a contention-oriented operating point, not as a representative deployment-wide multi-hop calibration.

Second, route confidence is a plausible abstraction, but evidence for its benefit comes primarily from the controlled route-conflict experiment. The random-topology matrix does not estimate its average benefit reliably because links are mostly high-PRR and cached paths are almost always one hop.

TABLE III

Controlled route-conflict fairness audit. Values are means over 100 paired seeds. The selected-route columns report the fraction of runs selecting the indicated route class.

Condition	MeshEcho PDR	No-confidence PDR	Δ PDR	Long selected	Short selected	Δ airtime (s)
No extra shadowing	1.000	1.000	0.000	0.000	1.000	0.000
Short branch +5 dB	0.980	0.980	0.000	0.020	0.980	0.000
Short branch +7.5 dB	0.999	0.784	+0.215	1.000	0.880	+15.148
Long branch +7.5 dB	1.000	1.000	0.000	0.000	1.000	0.000

Third, the best operating point depends on traffic. In repeated unicast, route caching is already strong and MeshEcho mainly improves efficiency. In mixed traffic, its reliability gain over source-route caching includes bounded recovery unavailable to the fixed baselines. This is a protocol-bundle result, not an isolated estimate of confidence ranking. Redundancy should therefore be treated as a per-context budget rather than a static network-wide mode.

Fourth, the two experiments answer different questions: route-conflict isolates admission under a known quality difference, while the random matrix measures end-to-end behavior. The main matrix is a common-harness comparison, not a fully budget-matched comparison of all recovery mechanisms; a mechanism study must disable recovery for all protocols or equalize retry and fallback budgets.

VII. Limitations and Future Work

This work is a packet-level simulation rather than a physical deployment. The model includes airtime, shadowing, probabilistic reception, half-duplex radios, collision, and capture, but not full waveforms, duty-cycle regulation, real firmware queues, or hardware sensitivity variation. The default 50-node, 3000 m, SF9 setting is a contention-oriented, high-PRR regime with almost no multi-hop cached paths. Repeated unicast and the unicast component of mixed traffic use eight fixed pairs, and no random-pair workload is included. Consequently, the current data are not sufficient to claim calibrated multi-hop deployment performance.

The comparison is fair at the common-harness level but not fully budget-matched: MeshEcho has bounded fallback, while the fixed baselines have no identical recovery controller. The frozen calm-mesh line uses a zero confidence-fallback threshold, while the exploratory Smart-CALM variants also allow timeout retries and profile changes. Their reliability cannot therefore be attributed to confidence ranking alone. Moreover, reception, forwarding jitter, and exploration share one stream in the frozen matrix: same seeds fix topology, traffic, and static shadowing, but not strict event-by-event common random numbers. The released independent-stream option removes this particular coupling for new reruns, but does not solve the different-transmission-count problem. All packet kinds also use the same configured payload length,

so source-route header growth and control-frame format differences are omitted from the airtime estimate.

The historical Smart-CALM ablation CSVs predate the no-fallback threshold fix included in the released implementation and are not used for the main quantitative claims. They remain in the repository only to preserve version traceability; a fresh ablation rerun is required before interpreting those variant rows.

The route-conflict ablation uses one seven-node topology, deterministic permitted edges, zero stochastic shadowing, staggered route-request timing, and a static channel. Its 100 seeds are repeated stochastic trials on that topology, not 100 independent network layouts. The confidence estimate and subsequent forwarding use the same static link realization, so stale or noisy confidence is not tested. Finally, the comparison does not include a separately implemented ETX or airtime-aware routing baseline.

Future work will calibrate a genuine multi-hop regime using a pre-registered direct-link PRR distribution and cached-route hop distribution, add random-pair workloads, equalize retry and fallback budgets, use the independent channel random stream (or keyed per-event draws), model variable packet headers, introduce time-varying links, and validate the design with hardware-in-the-loop experiments.

VIII. Conclusion

This paper presented a preliminary evaluation of MeshEcho confidence-aware routing for airtime-constrained LoRa mesh IoT networks. Using a reproducible packet-level simulator, we compared managed flooding, source-route caching, and the calm-mesh MeshEcho prototype. The results show that fixed routing modes have clear operating boundaries: flooding is reliable but airtime-intensive, while cached source routing is efficient but can lose unicast reliability in mixed traffic. In the frozen 50-node mixed matrix, MeshEcho reaches 0.735 ACK-confirmed unicast PDR with 840 s of airtime, compared with 0.389 PDR and 971 s for source-route caching. The controlled audit further shows that the route-ranking gain disappears under null, moderate, and reverse controls, which bounds the mechanism claim to conditions where link-quality evidence is informative. Because fixed pairs and unequal recovery budgets remain, these results should

be read as preliminary common-harness evidence rather than a universal ranking of LoRa mesh protocols.

- [16] MeshCore Project, “Meshcore,” <https://github.com/ripplebiz/MeshCore>, 2026, accessed 2026-05-11.

Reproducibility and Declarations

The MeshEcho simulator, routing implementation, experiment scripts, raw results, and paper source are released in the accompanying repository: <https://github.com/BH4ME/lora-mesh-routing-sim>. No human participants or personal data were involved. The authors report no conflicts of interest and no external funding. Author contributions cover protocol design, implementation, experiment design, analysis, and manuscript preparation. AI-based tools assisted with code inspection and language editing; the authors verified the implementation, numerical results, and final manuscript.

References

- [1] A. Augustin, J. Yi, T. Clausen, and W. M. Townsley, “A study of LoRa: Long range & low power networks for the internet of things,” *Sensors*, vol. 16, no. 9, p. 1466, 2016.
- [2] A. W.-L. Wong, S. L. Goh, M. K. Hasan, and S. Fattah, “Multi-hop and mesh for LoRa networks: Recent advancements, issues, and recommended applications,” *ACM Computing Surveys*, vol. 56, no. 6, pp. 1–43, 2024.
- [3] J. R. Cotrim and J. H. Kleinschmidt, “LoRaWAN mesh networks: A review and classification of multihop communication,” *Sensors*, vol. 20, no. 15, p. 4273, 2020.
- [4] A. Osorio, M. Calle, J. D. Soto, and J. E. Candelo-Becerra, “Routing in LoRaWAN: Overview and challenges,” *IEEE Communications Magazine*, vol. 58, no. 6, pp. 72–76, 2020.
- [5] R. P. Centelles, F. Freitag, R. Meseguer, and L. Navarro, “Beyond the star of stars: An introduction to multihop and mesh for LoRa and LoRaWAN,” *IEEE Pervasive Computing*, vol. 20, no. 2, pp. 63–72, 2021.
- [6] R. Pueyo Centelles, R. Meseguer, F. Freitag, R. Baig Viñas, and L. Navarro, “A minimalistic distance-vector routing protocol for LoRa mesh networks,” *IEEE Access*, vol. 12, pp. 128 941–128 962, 2024.
- [7] J. M. Sole, R. Pueyo Centelles, F. Freitag, and R. Meseguer, “Implementation of a LoRa mesh library,” *IEEE Access*, vol. 10, pp. 113 158–113 171, 2022.
- [8] G. Zhu, C.-H. Liao, T. Sakdejayont, I.-W. Lai, Y. Narusue, and H. Morikawa, “Improving the capacity of a mesh LoRa network by spreading-factor-based network clustering,” *IEEE Access*, vol. 7, pp. 21 584–21 596, 2019.
- [9] A.-U.-H. Ahmar, W. Joosen, and D. Hughes, “Smart-hop: Low-latency multi-hop networking for LoRa,” in *2022 18th International Conference on Distributed Computing in Sensor Systems (DCOSS)*, 2022, pp. 17–20.
- [10] L. Leonardi, L. Lo Bello, and G. Patti, “MRT-LoRa: A multi-hop real-time communication protocol for industrial IoT applications over LoRa networks,” *Computer Communications*, vol. 199, pp. 72–86, 2023.
- [11] T. Winter, P. Thubert, A. Brandt, J. Hui, R. Kelsey, P. Levis, K. Pister, R. Struik, J. P. Vasseur, and R. Alexander, “RPL: IPv6 routing protocol for low-power and lossy networks,” *IETF, RFC 6550*, 2012.
- [12] D. S. J. De Couto, D. Aguayo, J. Bicket, and R. Morris, “A high-throughput path metric for multi-hop wireless routing,” in *Proceedings of the 9th Annual International Conference on Mobile Computing and Networking*, ser. *MobiCom '03*, 2003, pp. 134–146.
- [13] Semtech Corporation, SX1261/2 Long Range, Low Power, sub-GHz RF Transceiver Datasheet, 2024, datasheet.
- [14] —, AN1200.13: LoRa Modem Designer’s Guide, 2013, application note.
- [15] Meshtastic Project, “Meshtastic documentation,” <https://meshtastic.org/docs/>, 2026, accessed 2026-05-11.